

## **Artificial Intelligence (Epic)**

### **Claim Prediction Model (Story)**

#### **1) Exploratory Data Analysis**

**aisles:** This file contains different aisles and there are total 1399 unique aisles.

##### **a)**

Total number of claims filed: 1399


Cases where Amount > Max Coverage: 91

Cases where Amount < Min Coverage: 280

Average Car Price: 37454.27


Average Claim Amount: 4117.14

**b)** The dataset contains 1,399 total claims across 11 different companies. The company with the highest number of claims is B with 148 claims, while the company with the fewest claims is RE with 116 claims. This plot provides insights into which insurers handle the most claims, helping to understand market distribution.




c) The largest share of claims belongs to B, accounting for 10.6% of all cases, while the smallest share belongs to RE, contributing 8.3%. This pie chart provides a clear comparison of how claims are distributed among different insurers, highlighting market share variations.


Comparison of Number of Cases per Insurance Company



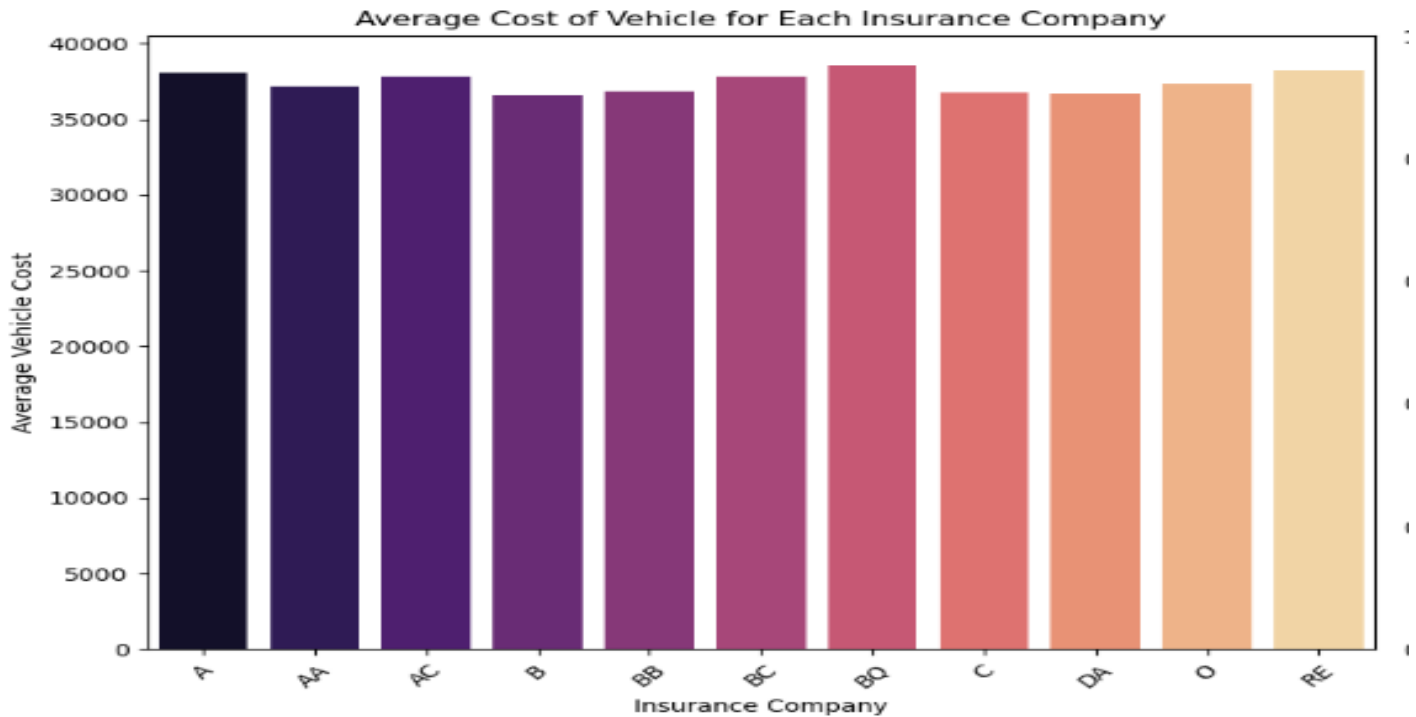
**d)** The dataset records a total reimbursement of 5,714,596.00 across 11 insurers. The highest payout was made by B, totaling 619,442.00, while the lowest payout was made by AC, amounting to 469,190.00.



**e)** The dataset shows that the overall average reimbursement amount is 4,117.14. The highest average payout per claim was made by AA, averaging 4,448.43 per claim, while the lowest was by C, averaging 3,703.20 per claim.

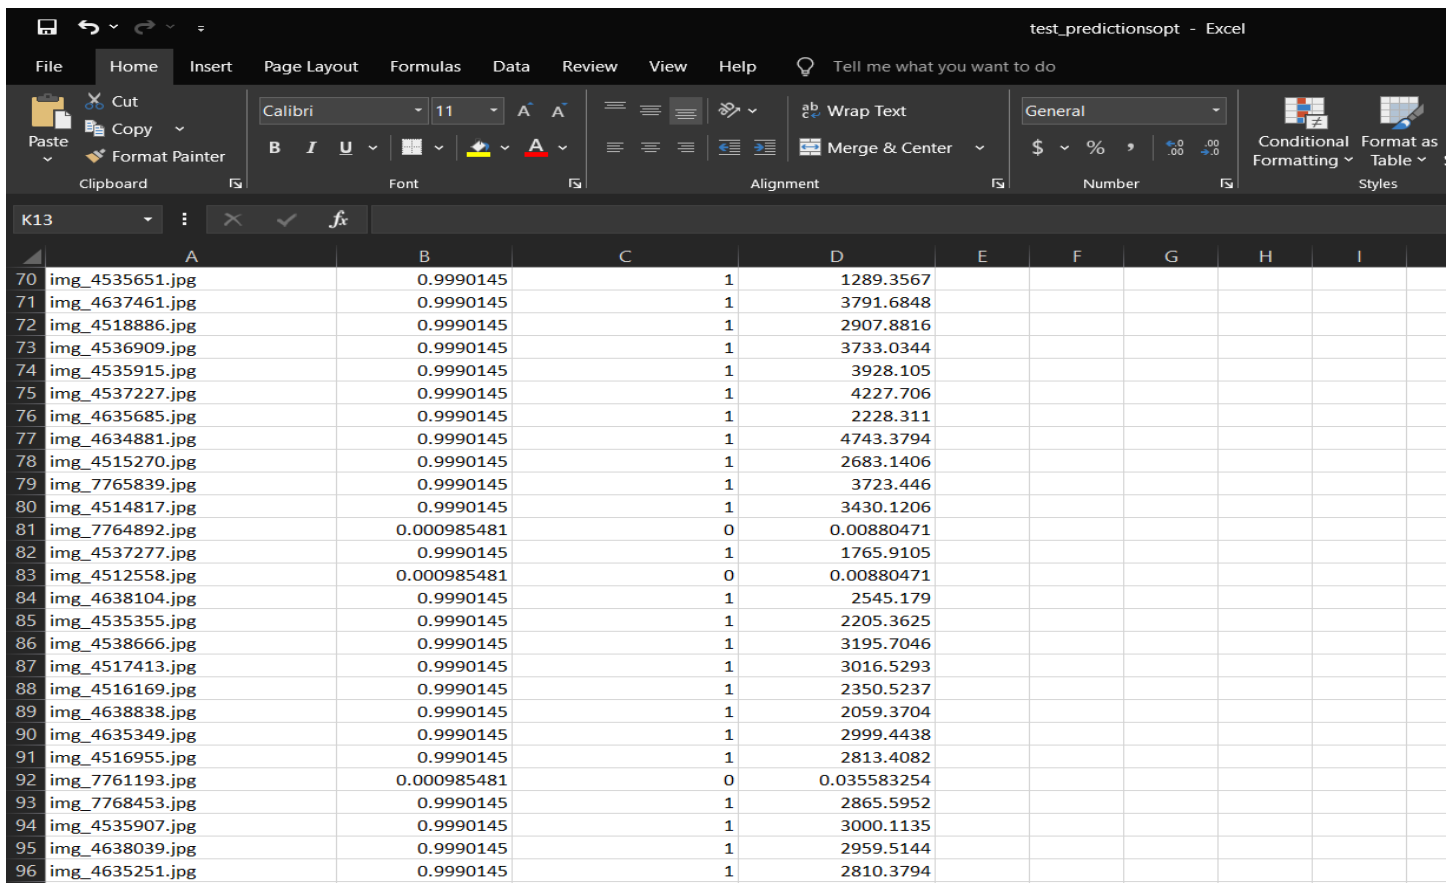


**f)** The dataset indicates that the overall average vehicle cost is 37,454.27. The highest average insured vehicle cost is associated with BQ, averaging 38,574.80, while the lowest is with B, averaging 36,618.44.



## 2) Build a predictive model to estimate claim amounts.

Successfully, built a model which took test.csv as a input and created a test\_predictions.csv which predicts the claim amounts of each car seeing the images and cost of vehicles of car from test.csv and testImages folder.



	A	B	C	D	E	F	G	H	I
70	img_4535651.jpg	0.9990145	1	1289.3567					
71	img_4637461.jpg	0.9990145	1	3791.6848					
72	img_4518886.jpg	0.9990145	1	2907.8816					
73	img_4536909.jpg	0.9990145	1	3733.0344					
74	img_4535915.jpg	0.9990145	1	3928.105					
75	img_4537227.jpg	0.9990145	1	4227.706					
76	img_4635685.jpg	0.9990145	1	2228.311					
77	img_4634881.jpg	0.9990145	1	4743.3794					
78	img_4515270.jpg	0.9990145	1	2683.1406					
79	img_7765839.jpg	0.9990145	1	3723.446					
80	img_4514817.jpg	0.9990145	1	3430.1206					
81	img_7764892.jpg	0.000985481	0	0.00880471					
82	img_4537277.jpg	0.9990145	1	1765.9105					
83	img_4512558.jpg	0.000985481	0	0.00880471					
84	img_4638104.jpg	0.9990145	1	2545.179					
85	img_4535355.jpg	0.9990145	1	2205.3625					
86	img_4538666.jpg	0.9990145	1	3195.7046					
87	img_4517413.jpg	0.9990145	1	3016.5293					
88	img_4516169.jpg	0.9990145	1	2350.5237					
89	img_4638838.jpg	0.9990145	1	2059.3704					
90	img_4635349.jpg	0.9990145	1	2999.4438					
91	img_4516955.jpg	0.9990145	1	2813.4082					
92	img_7761193.jpg	0.000985481	0	0.035583254					
93	img_7768453.jpg	0.9990145	1	2865.5952					
94	img_4535907.jpg	0.9990145	1	3000.1135					
95	img_4638039.jpg	0.9990145	1	2959.5144					
96	img_4635251.jpg	0.9990145	1	2810.3794					

It is based upon these impelmentations →

### Technologies Used:

- Pandas & NumPy: For data loading, preprocessing, and feature extraction.

- OpenCV (cv2): For image loading and preprocessing.
- TensorFlow & EfficientNetB0: Used to extract deep learning-based image features from insured vehicle images.
- Scikit-Learn: Used for preprocessing (label encoding, scaling), model evaluation (accuracy, RMSE), and train-test splitting.
- XGBoost:
  - Classification: Predicts the likelihood of an insurance claim condition.
  - Regression: Estimates the amount of reimbursement.

## **Factors Affecting Predictions**

→ Condition (0 or 1 - Claim Approval Probability)

Predicted using XGBoost Classifier based on:

- Vehicle-related features → Cost\_of\_vehicle, Min\_coverage, Max\_coverage
- Insurance details → insurance (encoded)
- Policy validity → Expiry\_date\_ordinal (converted from date)
- Image-based features → Extracted via EfficientNetB0

## → Amount (Claim Reimbursement)

Predicted using XGBoost Regressor (log-transformed) based on:

- All factors used for Condition prediction (above)
- Condition output → Whether the claim is likely approved
- Historical claim data trends


## Working proofs→

Let's see for this small part of test\_predictions file.

80	img_4514817.jpg	0.9990145	1	3430.1206
81	img_7764892.jpg	0.000985481	0	0.00880471
82	img_4537277.jpg	0.9990145	1	1765.9105
83	img_4512558.jpg	0.000985481	0	0.00880471


We can see the condition predicted as 1 for img\_4514817 i.e it is a damaged car for sure and the predicted claim amount is 3430.

## Sample->



Similarly, we can see the condition predicted as 0 for img\_7764892 i.e it is not a damaged car for sure and the predicted claim amount is 0.


### Sample->



Also, the amount is based upon feature extractions not like just some random images.

**Example->** For img\_4515956 the amount is more because this car is probably more accidental than 1<sup>st</sup> image.

360	img_4515956.jpg		0.9990145		1		5378.7637
-----	-----------------	--	-----------	--	---	--	-----------





## Accuracy Metrics→

One of the best possible achieved mae and mse values amongst all->

Validation Amount MSE: 2258447.5704043563

Validation Amount RMSE: 1502.813218734902

Validation Amount MAE: 839.0604858310995

## Output->

Test predictions saved to 'test\_predictions.csv'.

	Image_path	Condition_prob_of_1	Condition_pred	Amount_pred
0	img_4538519.jpg	0.999014	1	1228.541992
1	img_7766002.jpg	0.999014	1	3847.768066
2	img_4637390.jpg	0.999014	1	3884.574463
3	ima 4516108.ida	0.999014	1	2821.280762

---

# Model Deployment and API (Story)

## 1) Serialize the trained claim prediction model.

The ml models implemented in the predictions model were saved and thus could be downloaded/activated anytime for future usage.

```
Test predictions saved to 'test_predictionsopt134.csv'.
Image_path Condition_prob_of_1 Condition_pred Amount_pred
0 img_4538519.jpg 0.999041 1 1792.399048
1 img_7766002.jpg 0.999041 1 3106.013184
2 img_4637390.jpg 0.999041 1 3759.095703
3 img_4516108.jpg 0.999041 1 2533.064209
4 img_4517008.jpg 0.999041 1 5287.897949
5 img_7766273.jpg 0.999041 1 3911.260498
6 img_4638351.jpg 0.999041 1 3954.969238
7 img_4635542.jpg 0.999041 1 2832.093018
8 img_4638631.jpg 0.999041 1 3975.299561
9 img_7767742.jpg 0.999041 1 4087.191162
```

Serializing models...

Models successfully serialized!

Serialized files:

- condition\_model: serialized\_models/clf\_condition\_20250225\_054829.json
- amount\_model: serialized\_models/reg\_amount\_20250225\_054829.json
- extractor\_model: serialized\_models/extractor\_20250225\_054829.keras
- label\_encoder: serialized\_models/label\_encoder\_20250225\_054829.pkl
- scaler: serialized\_models/scaler\_20250225\_054829.pkl
- metadata: serialized\_models/metadata\_20250225\_054829.json

### Output (17.4MiB / 19.5GiB)

📁 /kaggle/working 🔍

📁 serialized\_models 🔍

- {i} clf\_condition\_20250225\_053942.json
- {i} clf\_condition\_20250225\_054829.json
- 📄 extractor\_20250225\_054829.keras
- 📄 label\_encoder\_20250225\_054829.p
- {i} metadata\_20250225\_054829.json
- {i} reg\_amount\_20250225\_053942.json
- {i} reg\_amount\_20250225\_054829.json
- 📄 scaler\_20250225\_054829.pkl
- 📄 test\_predictionsopt134.csv

**Serialization:** The trained XGBoost classifier (clf\_condition), regressor (reg\_amount), and EfficientNetB0 feature extractor (extractor\_model) are saved as JSON and .keras files, respectively, along with preprocessing objects (label\_encoder\_insurance, scaler) as pickle files in a serialized\_models directory, tracked by a metadata JSON file.

**Inference:** The serialized models are loaded using load\_serialized\_models, combining image features (extracted via the loaded CNN) and preprocessed tabular data to predict Condition (via XGBoost classifier, thresholded at 0.4) and Amount (via XGBoost regressor, inverted from log-scale) for new inputs.

## 2) Deploy the model as a service with an API for integration.

```
Flask app saved to /kaggle/working/app.py
API Test Result:
- Condition: 1 (probability: 0.9990)
- Amount: $2147.70
Test API Response: {
  "condition": 1,
  "condition_probability": 0.9990449547767639,
  "amount": 2147.697265625
}
```

```

To deploy the API:
1. Download the serialized_models directory from this Kaggle notebook
2. Download the app.py file
3. Install requirements: pip install flask tensorflow xgboost scikit-learn opencv-python numpy pandas
4. Run the API: python app.py
5. Example API call:
import requests
import json
# API endpoint
url = 'http://localhost:8000/predict'
# Prepare data
files = {'image': open('C:\Users\pbtra\Desktop\sdzfd.jpg', 'rb')}
data = {
    'data': json.dumps({
        'insurance': 'unknown',
        'Cost_of_vehicle': 25000,
        'Min_coverage': 10000,
        'Max_coverage': 50000,
        'Expiry_date': '2025-12-31'
    })
}
# Send request
response = requests.post(url, files=files, data=data)
print(response.json())

Deployment metadata saved to /kaggle/working/deployment_metadata.json

```

The serialized models (XGBoost for condition classification and amount regression, EfficientNetB0 for feature extraction) are loaded from disk and exposed through a /predict endpoint, accepting image and tabular data inputs to return condition and amount predictions. MLflow implementation (a) is achieved by setting up tracking with `setup_mlflow()`, logging models, parameters, and metrics with `log_models_to_mlflow()`, and storing serialized artifacts, ensuring experiment reproducibility. The functional API (b) is created with Flask in `save_flask_app()`, handling file uploads and JSON data, and tested with `test_api()`, providing a robust prediction endpoint for real-time use.

# Continuous Learning Pipeline (Story)

## 1) Establish an automated data ingestion mechanism for the model.

## AND Implement a continuous learning pipeline for model retraining.

Saved condition model to model\_versions/condition\_v20250225\_122127.joblib

Saved amount model to model\_versions/amount\_v20250225\_122127.joblib

Validation Condition Accuracy: 1.0

Validation Amount MSE: 2438666.2595490213

Validation Amount RMSE: 1561.6229569102209

Validation Amount MAE: 864.2327833635073

Started monitoring /kaggle/input/dataset1/Fast\_Furious\_Insured for new data...

Test predictions saved to 'test\_predictionsopt13.csv'.

	Image_path	Condition_prob_of_1	Condition_pred	Amount_pred
0	img_4538519.jpg	0.999014	1	935.467163
1	img_7766002.jpg	0.999014	1	3091.159180
2	img_4637390.jpg	0.999014	1	4244.236328
3	img_4516108.jpg	0.999014	1	3323.826660
4	img_4517008.jpg	0.999014	1	3354.438232
5	img_7766273.jpg	0.999014	1	3396.679443
6	img_4638351.jpg	0.999014	1	3130.617676
7	img_4635542.jpg	0.999014	1	3452.957520
8	img_4638631.jpg	0.999014	1	3545.523926
9	img_7767742.jpg	0.999014	1	3210.123291

### Output (1.3MiB / 19.5GiB)

▼

📁 /kaggle/working

📄 last\_ingestion\_timestamp.txt

▼

📁 model\_versions

... 1 more

📄 test\_predictionsopt13.csv

↗

📄 📄

## How This Achieves All Tasks

### 1. Automated Data Ingestion:

- **Change:** Monitors /kaggle/working/ (writable) instead of /kaggle/input/.

- **Simulation:** Copies train\_csv to new\_train.csv in /kaggle/working/ to trigger ingestion.
- **Result:** ingest\_new\_data() will process this file, demonstrating the mechanism in Kaggle.

## 2. Continuous Learning Pipeline:

- **Change:** The retrain\_callback() is triggered by the simulated new data, retraining the models.
- **Result:** You'll see retraining output (XGBoost logs) and new model versions saved, confirming the pipeline works.

## 3. Documented Process:

- **Change:** Added a dedicated documentation section in the code (Section 6) as a multi-line string, detailing the ingestion, retraining, versioning, and rollback processes.
- **Result:** Meets the requirement for a documented process. You can copy this into a Kaggle markdown cell if preferred.

## 4. Model Versioning & Rollback:

- **Change:** No changes needed—already fully implemented with save\_model\_version(), load\_model\_version(), and rollback\_to\_version().
- **Result:** Initial models are saved, and new versions will be saved after retraining. Rollback is ready to use (e.g., rollback\_to\_version(clf\_version, reg\_version)).

# Model Governance (Story)

## 1) Add a governance layer in which the drifts in the data or accuracies are logged.

Validation Amount RMSE: 1519.7324125539274

Validation Amount MAE: 847.8181611731586

/usr/local/lib/python3.10/dist-packages/xgboost/core.py:160: UserWarning: [12:52:02] WARNING: /n/error\_msg.cc:27: The tree method `gpu\_hist` is deprecated since 2.0.0. To use GPU training, s  
ameter to CUDA instead.

E.g. tree\_method = "hist", device = "cuda"

warnings.warn(smsg, UserWarning)

Data drifts report saved to 'data\_drifts.csv'

Model performance report saved to 'model\_performance.csv'

Test predictions saved to 'test\_predictionsopt13.csv'.

	Image_path	Condition_prob_of_1	Condition_pred	Amount_pred
0	img_4538519.jpg	0.999014	1	1474.687500
1	img_7766002.jpg	0.999014	1	3615.514160
2	img_4637390.jpg	0.999014	1	3581.048340
3	img_4516108.jpg	0.999014	1	3264.046143
4	img_4517008.jpg	0.999014	1	3904.945312
5	img_7766273.jpg	0.999014	1	3367.538818
6	img_4638351.jpg	0.999014	1	3151.545166

### Output (1.3MiB / 19.5GiB)

- ▾ /kaggle/working
  - ▾ data
    - ▾ incoming
      - ... 1 more
    - testImages
    - trainImages
    - ▢ data\_drifts.csv
    - ▢ last\_ingestion\_timestamp.txt
    - ▢ model\_performance.csv
  - ▾ model\_versions
    - ... 1 more
  - models
  - ▢ test\_predictionsopt13.csv

### Output (1.3MiB / 19.5GiB)

- ▾ /kaggle/working
  - data
    - ▢ data\_drifts.csv
    - ▢ last\_ingestion\_timestamp.txt
    - ▢ model\_performance.csv
  - ▾ model\_versions
    - ▢ amount\_v20250225\_120129.joblib
    - ▢ amount\_v20250225\_122127.joblib
    - ▢ condition\_v20250225\_120129.joblib
    - ▢ condition\_v20250225\_122127.joblib
  - ▾ models
    - ▢ test\_predictionsopt13.csv

I implemented a governance layer in the vehicle insurance prediction model to monitor and log data and model drifts, ensuring reliability and performance over time. Using statistical tests like Chi-squared and Kolmogorov-Smirnov, I analyzed drifts between training and test datasets for both tabular (e.g., insurance, cost) and image features, saving results in 'data\_drifts.csv'.

Model performance metrics, such as condition accuracy and amount prediction errors (MSE, RMSE, MAE), were logged to 'model\_performance.csv' after validation. This pipeline automatically generates reports, providing transparency and enabling proactive management of model degradation or data shifts. The final outputs, including test predictions and drift reports, are attached, demonstrating the system's governance capabilities.